

# 1\_Notebook\_Project\_LifeExpectancy\_LinearRegression

December 20, 2020

## 1 Linear Regression

1. Convert Business Problem to Data Science Problem
2. Load Data
3. Understand the Data
4. Data Preprocessing
5. Exploratory Data Analysis
6. Model Building
7. Model Diagnostics
8. Predictions and Evaluations

```
[1]: from google.colab import drive
drive.mount('/content/drive')
```

Mounted at /content/drive

### 1.1 Validate your System Libraries

Validate your System Libraries and if version is not updated, please update it.

```
[2]: # python version
# python --version

# python version (method 2)
from platform import python_version

print('python: {}'.format(python_version()))

#numpy version
import numpy as np
print('numpy: {}'.format(np.version.version))

#pandas version
import pandas as pd
print('pandas: {}'.format(pd.__version__))

#seaborn version
import seaborn as sns
```

```

print('seaborn: {}'.format(sns.__version__))

# matplotlib version
import matplotlib
print('matplotlib: {}'.format(matplotlib.__version__))

# sklearn version
import sklearn
print('The scikit-learn version is {}'.format(sklearn.__version__))

# statsmodels version
import statsmodels
print('statsmodels: {}'.format(statsmodels.__version__))

# statsmodels version
import imblearn
print('imblearn : {}'.format(imblearn.__version__))

# Pandas also provides a utility function, pd.show_versions(), which reports
→ the version of its dependencies as well:
# pd.show_versions(as_json=False)

```

```

python: 3.6.9
numpy: 1.19.4
pandas: 1.1.5
seaborn: 0.11.0
matplotlib: 3.2.2
The scikit-learn version is 0.22.2.post1.
statsmodels: 0.10.2
imblearn : 0.4.3

```

```

/usr/local/lib/python3.6/dist-packages/sklearn/externals/six.py:31:
FutureWarning: The module is deprecated in version 0.21 and will be removed in
version 0.23 since we've dropped support for Python 2.7. Please rely on the
official version of six (https://pypi.org/project/six/).
  "(https://pypi.org/project/six/).", FutureWarning)
/usr/local/lib/python3.6/dist-packages/sklearn/utils/deprecation.py:144:
FutureWarning: The sklearn.neighbors.base module is deprecated in version 0.22
and will be removed in version 0.24. The corresponding classes / functions
should instead be imported from sklearn.neighbors. Anything that cannot be
imported from sklearn.neighbors is now part of the private API.
  warnings.warn(message, FutureWarning)

```

## 1.2 1. Import Libraries

```
[3]: pd.set_option('display.max_rows', 800)
pd.set_option('display.max_columns', 500)

import matplotlib.pyplot as plt
%matplotlib inline

# import all libraries and dependencies for machine learning
from sklearn import preprocessing
from sklearn.model_selection import train_test_split
import statsmodels.api as sm
from sklearn.feature_selection import RFE
from sklearn.linear_model import LinearRegression
from statsmodels.stats.outliers_influence import variance_inflation_factor
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
import random
```

```
/usr/local/lib/python3.6/dist-packages/statsmodels/tools/_testing.py:19:
FutureWarning: pandas.util.testing is deprecated. Use the functions in the
public API at pandas.testing instead.
import pandas.util.testing as tm
```

## 1.3 2. Load Data

```
[4]: # Loading the dataset
df = pd.read_csv("/content/drive/MyDrive/Colab Notebooks/Life Expectancy Data.
→csv")
```

## 1.4 3. Understanding the data

```
[5]: #Pandas dataframe.info() function is used to get a concise summary of the
→dataframe.
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2938 entries, 0 to 2937
Data columns (total 22 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Country               2938 non-null   object
1   Year                  2938 non-null   int64
2   Status                2938 non-null   object
3   Life expectancy       2928 non-null   float64
4   Adult Mortality       2928 non-null   float64
5   infant deaths         2938 non-null   int64
6   Alcohol               2744 non-null   float64
```

7	percentage expenditure	2938 non-null	float64
8	Hepatitis B	2385 non-null	float64
9	Measles	2938 non-null	int64
10	BMI	2904 non-null	float64
11	under-five deaths	2938 non-null	int64
12	Polio	2919 non-null	float64
13	Total expenditure	2712 non-null	float64
14	Diphtheria	2919 non-null	float64
15	HIV/AIDS	2938 non-null	float64
16	GDP	2490 non-null	float64
17	Population	2286 non-null	float64
18	thinness 1-19 years	2904 non-null	float64
19	thinness 5-9 years	2904 non-null	float64
20	Income composition of resources	2771 non-null	float64
21	Schooling	2775 non-null	float64

dtypes: float64(16), int64(4), object(2)  
memory usage: 505.1+ KB

```
[6]: # Exploring the dataset
df.shape
```

```
[6]: (2938, 22)
```

## 2 Missing values

```
[7]: # Checking the null values in the dataset
df.isnull().sum()
```

```
[7]: Country          0
Year                0
Status              0
Life expectancy     10
Adult Mortality     10
infant deaths       0
Alcohol             194
percentage expenditure  0
Hepatitis B        553
Measles             0
BMI                 34
under-five deaths   0
Polio               19
Total expenditure   226
Diphtheria          19
HIV/AIDS            0
GDP                 448
Population          652
thinness 1-19 years  34
```

```

    thinness 5-9 years          34
Income composition of resources 167
Schooling                     163
dtype: int64

```

[8]: *#The describe() function computes a summary of statistics pertaining to the DataFrame columns.*  
df.describe()

```

[8]:
      Year  Life expectancy  Adult Mortality  infant deaths  \
count  2938.000000      2928.000000      2928.000000      2938.000000
mean    2007.518720       69.224932       164.796448       30.303948
std       4.613841       9.523867       124.292079      117.926501
min     2000.000000       36.300000        1.000000        0.000000
25%     2004.000000       63.100000       74.000000        0.000000
50%     2008.000000       72.100000      144.000000        3.000000
75%     2012.000000       75.700000      228.000000       22.000000
max     2015.000000       89.000000      723.000000     1800.000000

      Alcohol  percentage expenditure  Hepatitis B      Measles  \
count  2744.000000      2938.000000  2385.000000      2938.000000
mean     4.602861       738.251295    80.940461     2419.592240
std     4.052413     1987.914858    25.070016    11467.272489
min     0.010000        0.000000     1.000000        0.000000
25%     0.877500        4.685343    77.000000        0.000000
50%     3.755000       64.912906    92.000000       17.000000
75%     7.702500      441.534144    97.000000      360.250000
max     17.870000     19479.911610    99.000000    212183.000000

      BMI  under-five deaths      Polio  Total expenditure  \
count  2904.000000      2938.000000  2919.000000      2712.000000
mean     38.321247      42.035739    82.550188        5.93819
std     20.044034     160.445548    23.428046        2.49832
min      1.000000        0.000000     3.000000        0.37000
25%     19.300000        0.000000    78.000000        4.26000
50%     43.500000        4.000000    93.000000        5.75500
75%     56.200000       28.000000    97.000000        7.49250
max     87.300000     2500.000000    99.000000       17.60000

      Diphtheria      HIV/AIDS      GDP      Population  \
count  2919.000000  2938.000000  2490.000000  2.286000e+03
mean     82.324084    1.742103   7483.158469  1.275338e+07
std     23.716912    5.077785  14270.169342  6.101210e+07
min      2.000000    0.100000    1.681350   3.400000e+01
25%     78.000000    0.100000   463.935626  1.957932e+05
50%     93.000000    0.100000  1766.947595  1.386542e+06
75%     97.000000    0.800000  5910.806335  7.420359e+06
max     99.000000   50.600000  119172.741800  1.293859e+09

```

	thinness 1-19 years	thinness 5-9 years \
count	2904.000000	2904.000000
mean	4.839704	4.870317
std	4.420195	4.508882
min	0.100000	0.100000
25%	1.600000	1.500000
50%	3.300000	3.300000
75%	7.200000	7.200000
max	27.700000	28.600000

	Income composition of resources	Schooling
count	2771.000000	2775.000000
mean	0.627551	11.992793
std	0.210904	3.358920
min	0.000000	0.000000
25%	0.493000	10.100000
50%	0.677000	12.300000
75%	0.779000	14.300000
max	0.948000	20.700000

```
[9]: # print the 5 records of the dataset by default. Pass number how many record
      ↳you want look
df.head()
```

```
[9]: Country Year Status Life expectancy Adult Mortality \
0 Afghanistan 2015 Developing 65.0 263.0
1 Afghanistan 2014 Developing 59.9 271.0
2 Afghanistan 2013 Developing 59.9 268.0
3 Afghanistan 2012 Developing 59.5 272.0
4 Afghanistan 2011 Developing 59.2 275.0
```

	infant deaths	Alcohol percentage	expenditure	Hepatitis B	Measles \
0	62	0.01	71.279624	65.0	1154
1	64	0.01	73.523582	62.0	492
2	66	0.01	73.219243	64.0	430
3	69	0.01	78.184215	67.0	2787
4	71	0.01	7.097109	68.0	3013

	BMI	under-five deaths	Polio	Total expenditure	Diphtheria \
0	19.1	83	6.0	8.16	65.0
1	18.6	86	58.0	8.18	62.0
2	18.1	89	62.0	8.13	64.0
3	17.6	93	67.0	8.52	67.0
4	17.2	97	68.0	7.87	68.0

	HIV/AIDS	GDP	Population	thinness 1-19 years \
0	0.1	584.259210	33736494.0	17.2

1	0.1	612.696514	327582.0	17.5
2	0.1	631.744976	31731688.0	17.7
3	0.1	669.959000	3696958.0	17.9
4	0.1	63.537231	2978599.0	18.2

	thinness 5-9 years	Income composition of resources	Schooling
0	17.3	0.479	10.1
1	17.5	0.476	10.0
2	17.7	0.470	9.9
3	18.0	0.463	9.8
4	18.2	0.454	9.5

```
[10]: # print the last 5 records of the dataset by default. Pass number how many
      ↪ record you want look
      df.tail()
```

```
[10]:      Country  Year      Status  Life expectancy  Adult Mortality \
2933  Zimbabwe  2004  Developing           44.3           723.0
2934  Zimbabwe  2003  Developing           44.5           715.0
2935  Zimbabwe  2002  Developing           44.8            73.0
2936  Zimbabwe  2001  Developing           45.3          686.0
2937  Zimbabwe  2000  Developing           46.0          665.0
```

	infant deaths	Alcohol	percentage expenditure	Hepatitis B	Measles	\
2933	27	4.36	0.0	68.0	31	
2934	26	4.06	0.0	7.0	998	
2935	25	4.43	0.0	73.0	304	
2936	25	1.72	0.0	76.0	529	
2937	24	1.68	0.0	79.0	1483	

	BMI	under-five deaths	Polio	Total expenditure	Diphtheria	\
2933	27.1	42	67.0	7.13	65.0	
2934	26.7	41	7.0	6.52	68.0	
2935	26.3	40	73.0	6.53	71.0	
2936	25.9	39	76.0	6.16	75.0	
2937	25.5	39	78.0	7.10	78.0	

	HIV/AIDS	GDP	Population	thinness	1-19 years	\
2933	33.6	454.366654	12777511.0		9.4	
2934	36.7	453.351155	12633897.0		9.8	
2935	39.8	57.348340	125525.0		1.2	
2936	42.1	548.587312	12366165.0		1.6	
2937	43.5	547.358879	12222251.0		11.0	

	thinness 5-9 years	Income composition of resources	Schooling
2933	9.4	0.407	9.2
2934	9.9	0.418	9.5
2935	1.3	0.427	10.0

2936	1.7	0.427	9.8
2937	11.2	0.434	9.8

```
[11]: num_col = df.select_dtypes(include=np.number).columns
print("Numerical columns: \n",num_col)

cat_col = df.select_dtypes(exclude=np.number).columns
print("Categorical columns: \n",cat_col)
```

Numerical columns:

```
Index(['Year', 'Life expectancy ', 'Adult Mortality', 'infant deaths',
      'Alcohol', 'percentage expenditure', 'Hepatitis B', 'Measles ', ' BMI ',
      'under-five deaths ', 'Polio', 'Total expenditure', 'Diphtheria ',
      ' HIV/AIDS', 'GDP', 'Population', ' thinness 1-19 years',
      ' thinness 5-9 years', 'Income composition of resources', 'Schooling'],
      dtype='object')
```

Categorical columns:

```
Index(['Country', 'Status'], dtype='object')
```

## 2.1 4. Data Pre-processing

```
[12]: # Remove the extra space from column names
```

```
df = df.rename(columns=lambda x: x.strip())
```

```
[13]: # Import label encoder
```

```
from sklearn import preprocessing
```

```
# label_encoder object knows how to understand word labels.
```

```
label_encoder = preprocessing.LabelEncoder()
```

```
# Encode labels in column 'Status'.
```

```
df['Status'] = label_encoder.fit_transform(df['Status'])
```

```
df.head(10)
```

```
[13]:
```

	Country	Year	Status	Life expectancy	Adult Mortality	infant deaths	\
0	Afghanistan	2015	1	65.0	263.0	62	
1	Afghanistan	2014	1	59.9	271.0	64	
2	Afghanistan	2013	1	59.9	268.0	66	
3	Afghanistan	2012	1	59.5	272.0	69	
4	Afghanistan	2011	1	59.2	275.0	71	
5	Afghanistan	2010	1	58.8	279.0	74	
6	Afghanistan	2009	1	58.6	281.0	77	
7	Afghanistan	2008	1	58.1	287.0	80	
8	Afghanistan	2007	1	57.5	295.0	82	
9	Afghanistan	2006	1	57.3	295.0	84	



	Alcohol	percentage expenditure	Hepatitis B	Measles	BMI	\
0	0.01	71.279624	65.0	1154	19.1	
1	0.01	73.523582	62.0	492	18.6	
2	0.01	73.219243	64.0	430	18.1	
3	0.01	78.184215	67.0	2787	17.6	
4	0.01	7.097109	68.0	3013	17.2	
5	0.01	79.679367	66.0	1989	16.7	
6	0.01	56.762217	63.0	2861	16.2	
7	0.03	25.873925	64.0	1599	15.7	
8	0.02	10.910156	63.0	1141	15.2	
9	0.03	17.171518	64.0	1990	14.7	

	under-five deaths	Polio	Total expenditure	Diphtheria	HIV/AIDS	\
0	83	6.0	8.16	65.0	0.1	
1	86	58.0	8.18	62.0	0.1	
2	89	62.0	8.13	64.0	0.1	
3	93	67.0	8.52	67.0	0.1	
4	97	68.0	7.87	68.0	0.1	
5	102	66.0	9.20	66.0	0.1	
6	106	63.0	9.42	63.0	0.1	
7	110	64.0	8.33	64.0	0.1	
8	113	63.0	6.73	63.0	0.1	
9	116	58.0	7.43	58.0	0.1	

	GDP	Population	thinness	1-19 years	thinness 5-9 years	\
0	584.259210	33736494.0		17.2	17.3	
1	612.696514	327582.0		17.5	17.5	
2	631.744976	31731688.0		17.7	17.7	
3	669.959000	3696958.0		17.9	18.0	
4	63.537231	2978599.0		18.2	18.2	
5	553.328940	2883167.0		18.4	18.4	
6	445.893298	284331.0		18.6	18.7	
7	373.361116	2729431.0		18.8	18.9	
8	369.835796	26616792.0		19.0	19.1	
9	272.563770	2589345.0		19.2	19.3	

	Income composition of resources	Schooling
0	0.479	10.1
1	0.476	10.0
2	0.470	9.9
3	0.463	9.8
4	0.454	9.5
5	0.448	9.2
6	0.434	8.9
7	0.433	8.7
8	0.415	8.4
9	0.405	8.1

```
[14]: print(df.isna().sum())  
      print(df.shape)
```

```
Country          0  
Year             0  
Status           0  
Life expectancy  10  
Adult Mortality  10  
infant deaths    0  
Alcohol          194  
percentage expenditure  0  
Hepatitis B      553  
Measles          0  
BMI              34  
under-five deaths  0  
Polio            19  
Total expenditure 226  
Diphtheria       19  
HIV/AIDS         0  
GDP              448  
Population       652  
thinness 1-19 years  34  
thinness 5-9 years  34  
Income composition of resources 167  
Schooling        163  
dtype: int64  
(2938, 22)
```


```
[15]: # Replace using mean  
      for i in df.columns.drop('Country'):  
          df[i].fillna(df[i].mean(), inplace = True)
```

## 2.2 5. Exploratory Data Analysis

```
[16]: # Let's check the distribution of y variable (Life Expectancy)  
      plt.figure(figsize=(8,8), dpi= 80)  
      sns.boxplot(df['Life expectancy'])  
      plt.title('Life expectancy Box Plot')  
      plt.show()
```

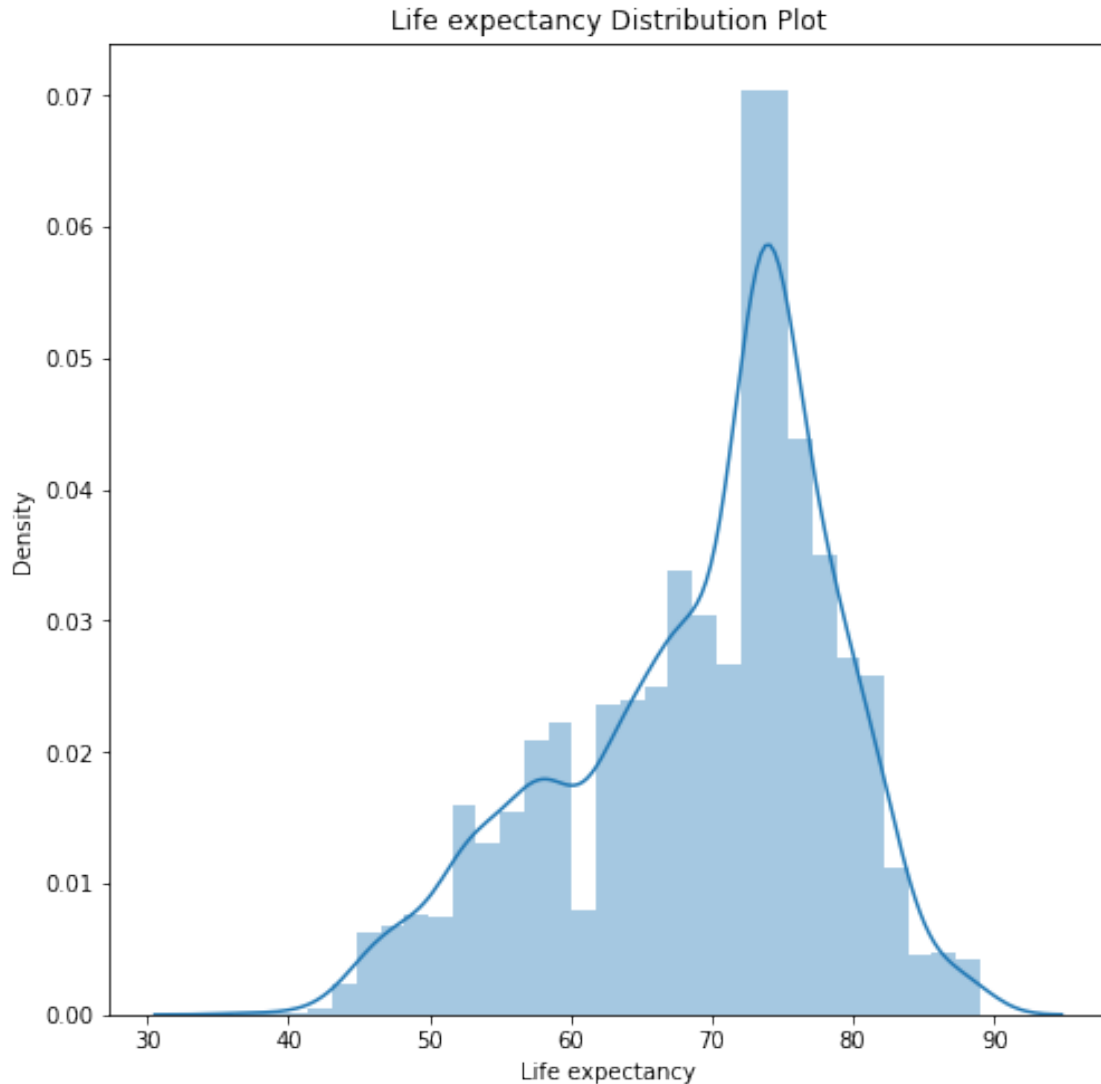
/usr/local/lib/python3.6/dist-packages/seaborn/\_decorators.py:43: FutureWarning:  
Pass the following variable as a keyword arg: x. From version 0.12, the only  
valid positional argument will be `data`, and passing other arguments without an  
explicit keyword will result in an error or misinterpretation.

FutureWarning



```
[17]: plt.figure(figsize=(8,8))  
plt.title('Life expectancy Distribution Plot')  
sns.distplot(df['Life expectancy']);
```

```
/usr/local/lib/python3.6/dist-packages/seaborn/distributions.py:2551:  
FutureWarning: `distplot` is a deprecated function and will be removed in a  
future version. Please adapt your code to use either `displot` (a figure-level  
function with similar flexibility) or `histplot` (an axes-level function for  
histograms).  
warnings.warn(msg, FutureWarning)
```



**Summary :** The y variable is having very few outliers and is almost linearly distributed. So the assumption for linear regression holds true

```
[18]: num_col = df.select_dtypes(include=np.number).columns
      print("Numerical columns: \n",num_col)

      cat_col = df.select_dtypes(exclude=np.number).columns
      print("Categorical columns: \n",cat_col)
```

Numerical columns:

```
Index(['Year', 'Status', 'Life expectancy', 'Adult Mortality', 'infant deaths',
      'Alcohol', 'percentage expenditure', 'Hepatitis B', 'Measles', 'BMI',
      'under-five deaths', 'Polio', 'Total expenditure', 'Diphtheria',
```

```

'HIV/AIDS', 'GDP', 'Population', 'thinness 1-19 years',
'thinness 5-9 years', 'Income composition of resources', 'Schooling'],
dtype='object')
Categorical columns:
Index(['Country'], dtype='object')

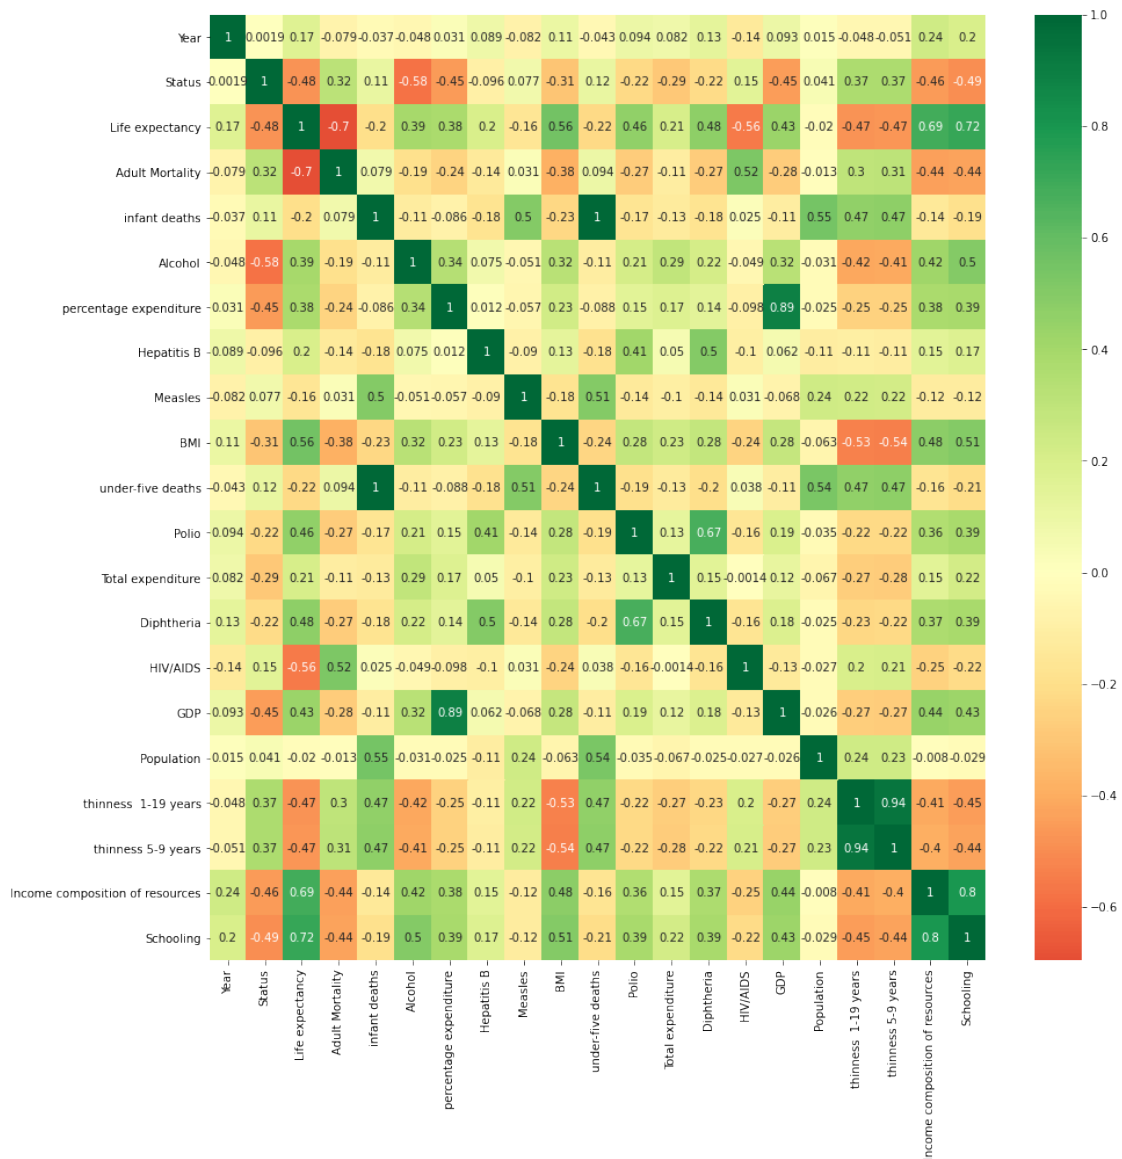
```

[19]: # Let's check the multicollinearity of features by checking the correlation  
→*matric*

```

plt.figure(figsize=(15,15))
p=sns.heatmap(df[num_col].corr(), annot=True,cmap='RdYlGn',center=0)

```



```
[20]: # Pair Plots to know the relation between different features
ax = sns.pairplot(df[num_col])
```

Output hidden; open in <https://colab.research.google.com> to view.

Few of the features are having the linear relationship with y variable. So linear regression would be good approach for the same

## 2.3 6. Model Building

```
[21]: # Train test split
X=df.drop(columns=['Life expectancy','Country'])
y=df[['Life expectancy']]

X_train, X_test, y_train, y_test = train_test_split(X,y,test_size=0.
→3,random_state=1234)
```

```
[22]: X.columns
```

```
[22]: Index(['Year', 'Status', 'Adult Mortality', 'infant deaths', 'Alcohol',
'percentage expenditure', 'Hepatitis B', 'Measles', 'BMI',
'under-five deaths', 'Polio', 'Total expenditure', 'Diphtheria',
'HIV/AIDS', 'GDP', 'Population', 'thinness 1-19 years',
'thinness 5-9 years', 'Income composition of resources', 'Schooling'],
dtype='object')
```

```
[23]: X.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2938 entries, 0 to 2937
Data columns (total 20 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Year                                  2938 non-null   int64
1   Status                               2938 non-null   int64
2   Adult Mortality                      2938 non-null   float64
3   infant deaths                        2938 non-null   int64
4   Alcohol                              2938 non-null   float64
5   percentage expenditure               2938 non-null   float64
6   Hepatitis B                          2938 non-null   float64
7   Measles                              2938 non-null   int64
8   BMI                                  2938 non-null   float64
9   under-five deaths                   2938 non-null   int64
10  Polio                                2938 non-null   float64
11  Total expenditure                   2938 non-null   float64
12  Diphtheria                          2938 non-null   float64
13  HIV/AIDS                            2938 non-null   float64
```

```

14 GDP                                2938 non-null    float64
15 Population                          2938 non-null    float64
16 thinness 1-19 years                 2938 non-null    float64
17 thinness 5-9 years                  2938 non-null    float64
18 Income composition of resources      2938 non-null    float64
19 Schooling                           2938 non-null    float64
dtypes: float64(15), int64(5)
memory usage: 459.2 KB

```

## 2.4 Approach 1 : Adding 1 variable after 1

### 2.4.1 Building model with 1 variable

```

[24]: # Select only one feature in regression model
X_train1 = X_train['Income composition of resources']

```

```

[25]: # Add a constant
X_train1 = sm.add_constant(X_train1)

# Create a first ols model
model_1 = sm.OLS(y_train, X_train1).fit()

```

```

[26]: # Check parameters created
model_1.params

```

```

[26]: const                48.440947
Income composition of resources  33.059741
dtype: float64

```

```

[27]: # Summary of the model
print(model_1.summary())

```

```

                                OLS Regression Results
=====
Dep. Variable:                Life expectancy    R-squared:                0.490
Model:                        OLS               Adj. R-squared:            0.490
Method:                        Least Squares      F-statistic:              1974.
Date:                          Sun, 20 Dec 2020    Prob (F-statistic):       1.09e-302
Time:                          16:09:50           Log-Likelihood:           -6894.3
No. Observations:              2056              AIC:                     1.379e+04
Df Residuals:                  2054              BIC:                     1.380e+04
Df Model:                      1
Covariance Type:               nonrobust
=====
=====
                                coef      std err          t      P>|t|
-----
[0.025      0.975]
-----
const                48.4409      0.492      98.412      0.000

```

```

47.476      49.406
Income composition of resources      33.0597      0.744      44.427      0.000
31.600      34.519
=====
Omnibus:                        138.959      Durbin-Watson:                        2.047
Prob(Omnibus):                  0.000      Jarque-Bera (JB):                    617.560
Skew:                          0.121      Prob(JB):                          7.92e-135
Kurtosis:                      5.674      Cond. No.                          6.86
=====

```

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

R-squared is very less for the model, so there's a need to add more features as R-square itself is not able to explain the expenses

## 2.4.2 Building model with 2 variable

```

[28]: # Add one more feature in regression model
X_train2 = X_train[['Income composition of resources', 'Schooling']]

```

```

[29]: # Add a constant
X_train2 = sm.add_constant(X_train2)

# Create second ols model
model_2 = sm.OLS(y_train, X_train2).fit()

```

```

[30]: # Check parameters created
model_2.params

```

```

[30]: const                43.145928
Income composition of resources  16.273079
Schooling                    1.320315
dtype: float64

```

```

[31]: # Summary of the model
print(model_2.summary())

```

```

                                OLS Regression Results
=====
Dep. Variable:      Life expectancy      R-squared:      0.562
Model:              OLS      Adj. R-squared:      0.561
Method:             Least Squares      F-statistic:      1316.
Date:              Sun, 20 Dec 2020      Prob (F-statistic):      0.00
Time:              16:09:50      Log-Likelihood:      -6738.3
No. Observations:      2056      AIC:      1.348e+04
Df Residuals:      2053      BIC:      1.350e+04
Df Model:              2

```



```

Covariance Type:          nonrobust
=====
=====
                                coef    std err          t      P>|t|
-----
[0.025    0.975]
-----
const                43.1459    0.540    79.895    0.000
42.087    44.205
Income composition of resources    16.2731    1.146    14.197    0.000
14.025    18.521
Schooling            1.3203    0.072    18.340    0.000
1.179    1.461
=====
Omnibus:            182.792    Durbin-Watson:           2.037
Prob(Omnibus):      0.000    Jarque-Bera (JB):        596.381
Skew:              -0.427    Prob(JB):               3.14e-130
Kurtosis:          5.497    Cond. No.                101.
=====

```

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

R-squared and Adj. R squared has increased for the model, but we can still improvise over it so let's add more features

### 2.4.3 Building model with 3 variable

```

[32]: # Adding one more feature in regression model
X_train3 = X_train[['Income composition of resources', 'Schooling', 'Adult_
↳Mortality']]

```

```

[33]: # Add a constant
X_train3 = sm.add_constant(X_train3)

# Create third fitted model
model_3 = sm.OLS(y_train, X_train3).fit()

```

```

[34]: # Check parameters created
model_3.params

```

```

[34]: const                56.227689
Income composition of resources    10.637516
Schooling                      1.003654
Adult Mortality                -0.034790
dtype: float64

```

```

[35]: # Summary of the model
print(model_3.summary())

```

```

                                OLS Regression Results
=====
Dep. Variable:                Life expectancy    R-squared:                0.721
Model:                        OLS              Adj. R-squared:           0.720
Method:                       Least Squares    F-statistic:             1765.
Date:                         Sun, 20 Dec 2020  Prob (F-statistic):    0.00
Time:                         16:09:51         Log-Likelihood:          -6275.3
No. Observations:             2056            AIC:                    1.256e+04
Df Residuals:                 2052            BIC:                    1.258e+04
Df Model:                     3
Covariance Type:              nonrobust
=====
=====
                                coef      std err          t      P>|t|
-----
[0.025      0.975]
-----
const                        56.2277      0.577      97.502      0.000
55.097      57.359
Income composition of resources  10.6375      0.930      11.438      0.000
8.814      12.461
Schooling                    1.0037      0.058      17.236      0.000
0.889      1.118
Adult Mortality              -0.0348      0.001     -34.168      0.000
-0.037      -0.033
=====
Omnibus:                     379.309    Durbin-Watson:           1.962
Prob(Omnibus):               0.000    Jarque-Bera (JB):       1628.478
Skew:                       -0.829    Prob(JB):               0.00
Kurtosis:                   7.032    Cond. No.               1.72e+03
=====

```

Warnings:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 1.72e+03. This might indicate that there are strong multicollinearity or other numerical problems.

We have achieved a R-squared of 0.72 by manually picking the highly correlated variables. Now lets use RFE to select the independent variables which accurately predicts the dependent variable Life expectancy.

## 2.5 Approach 2 : RFE and eleminating by using p-value and VIF

```

[36]: # Running RFE with important column count to be 15
lm = LinearRegression()
lm.fit(X_train, y_train)

```

```
rfe = RFE(lm, 15)
rfe = rfe.fit(X_train, y_train)
```

```
/usr/local/lib/python3.6/dist-packages/sklearn/utils/validation.py:760:
DataConversionWarning: A column-vector y was passed when a 1d array was
expected. Please change the shape of y to (n_samples, ), for example using
ravel().
y = column_or_1d(y, warn=True)
```

```
[37]: list(zip(X_train.columns, rfe.support_, rfe.ranking_))
```

```
[37]: [('Year', False, 2),
      ('Status', True, 1),
      ('Adult Mortality', True, 1),
      ('infant deaths', True, 1),
      ('Alcohol', True, 1),
      ('percentage expenditure', False, 3),
      ('Hepatitis B', True, 1),
      ('Measles', False, 5),
      ('BMI', True, 1),
      ('under-five deaths', True, 1),
      ('Polio', True, 1),
      ('Total expenditure', True, 1),
      ('Diphtheria', True, 1),
      ('HIV/AIDS', True, 1),
      ('GDP', False, 4),
      ('Population', False, 6),
      ('thinness 1-19 years', True, 1),
      ('thinness 5-9 years', True, 1),
      ('Income composition of resources', True, 1),
      ('Schooling', True, 1)]
```

```
[38]: # Selecting the important features (in the support)
imp_columns = X_train.columns[rfe.support_]
imp_columns
```

```
[38]: Index(['Status', 'Adult Mortality', 'infant deaths', 'Alcohol', 'Hepatitis B',
          'BMI', 'under-five deaths', 'Polio', 'Total expenditure', 'Diphtheria',
          'HIV/AIDS', 'thinness 1-19 years', 'thinness 5-9 years',
          'Income composition of resources', 'Schooling'],
          dtype='object')
```

```
[39]: # Creating X_train dataframe with RFE selected variables
X_train_rfe = X_train[imp_columns]
```

After passing the arbitrary selected columns by RFE we will manually evaluate each models p-value and VIF value. Unless we find the acceptable range for p-values and VIF we keep dropping the variables one at a time based on below criteria. - High p-value High VIF : Drop the variable - High p-value Low VIF : Drop the variable with high p-value first - Low p-value Low VIF : accept the variable

### 2.5.1 Checking VIF

Variance Inflation Factor or VIF, gives a basic quantitative idea about how much the feature variables are correlated with each other. It is an extremely important parameter to test our linear model. The formula for calculating VIF is:

```
[40]: random.seed(0)

# Add a constant
X_train_rfec = sm.add_constant(X_train_rfe)

# Build the model with RFE features
lm_rfe = sm.OLS(y_train,X_train_rfec).fit()

#Summary of linear model
print(lm_rfe.summary())
```

```

                        OLS Regression Results
=====
Dep. Variable:          Life expectancy    R-squared:                0.820
Model:                  OLS               Adj. R-squared:         0.819
Method:                 Least Squares     F-statistic:             620.1
Date:                  Sun, 20 Dec 2020   Prob (F-statistic):       0.00
Time:                  16:09:51          Log-Likelihood:          -5823.0
No. Observations:      2056             AIC:                    1.168e+04
Df Residuals:          2040             BIC:                    1.177e+04
Df Model:               15
Covariance Type:       nonrobust
=====
=====

```

		coef	std err	t	P> t
[0.025	0.975]				
-----					
const		56.3507	0.815	69.162	0.000
54.753	57.949				
Status		-2.0823	0.313	-6.663	0.000
-2.695	-1.469				
Adult Mortality		-0.0200	0.001	-20.474	0.000
-0.022	-0.018				
infant deaths		0.0944	0.010	9.677	0.000
0.075	0.114				
Alcohol		0.0403	0.031	1.289	0.198
-0.021	0.102				
Hepatitis B		-0.0209	0.005	-4.303	0.000
-0.030	-0.011				
BMI		0.0488	0.006	7.963	0.000
0.037	0.061				
under-five deaths		-0.0714	0.007	-9.983	0.000

-0.085	-0.057				
Polio		0.0272	0.005	4.962	0.000
0.016	0.038				
Total expenditure		0.0768	0.042	1.835	0.067
-0.005	0.159				
Diphtheria		0.0456	0.006	7.807	0.000
0.034	0.057				
HIV/AIDS		-0.4970	0.024	-20.592	0.000
-0.544	-0.450				
thinness 1-19 years		-0.0739	0.061	-1.212	0.226
-0.193	0.046				
thinness 5-9 years		0.0032	0.060	0.054	0.957
-0.114	0.120				
Income composition of resources		6.4836	0.769	8.435	0.000
4.976	7.991				
Schooling		0.6908	0.051	13.544	0.000
0.591	0.791				

Omnibus:	110.684	Durbin-Watson:	1.979
Prob(Omnibus):	0.000	Jarque-Bera (JB):	321.018
Skew:	-0.238	Prob(JB):	1.96e-70
Kurtosis:	4.876	Cond. No.	2.45e+03

Warnings:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 2.45e+03. This might indicate that there are strong multicollinearity or other numerical problems.

Since the p value for few of the features is not significant, we need to drop it but before that let's check the VIF score as well

```
[41]: # Create a dataframe that will contain the names of all the feature variables
      →and their respective VIFs
vif = pd.DataFrame()
vif['Features'] = X_train_rfe.columns
vif['VIF'] = [variance_inflation_factor(X_train_rfe.values, i) for i in
      →range(X_train_rfe.shape[1])]
vif['VIF'] = round(vif['VIF'], 2)
vif = vif.sort_values(by = "VIF", ascending = False)
vif
```

```
[41]:
```

	Features	VIF
6	under-five deaths	178.16
2	infant deaths	177.70
14	Schooling	44.59
13	Income composition of resources	30.42
9	Diphtheria	30.31

7		Polio	26.28
11	thinness 1-19 years		19.47
12	thinness 5-9 years		19.31
4	Hepatitis B		19.00
5	BMI		8.28
8	Total expenditure		7.74
0	Status		7.13
1	Adult Mortality		4.42
3	Alcohol		4.35
10	HIV/AIDS		1.70

Since the variable **thinness 5-9 years** is having a very high p value , we would remove the feature from training dataset

```
[42]: # Dropping insignificant variables

X_train_rfe1 = X_train_rfe.drop(['thinness 5-9 years'], 1,)

# Adding a constant variable and Build a second fitted model

X_train_rfe1c = sm.add_constant(X_train_rfe1)
lm_rfe1 = sm.OLS(y_train, X_train_rfe1c).fit()

#Summary of linear model
print(lm_rfe1.summary())
```

#### OLS Regression Results

```
=====
Dep. Variable:      Life expectancy      R-squared:      0.820
Model:              OLS                  Adj. R-squared: 0.819
Method:             Least Squares        F-statistic:    664.7
Date:               Sun, 20 Dec 2020      Prob (F-statistic): 0.00
Time:               16:09:51              Log-Likelihood: -5823.0
No. Observations:   2056                  AIC:            1.168e+04
Df Residuals:       2041                  BIC:            1.176e+04
Df Model:           14
Covariance Type:    nonrobust
=====
```

```
=====
                                coef      std err          t      P>|t|
-----
[0.025      0.975]
-----
const                                56.3537      0.813      69.336      0.000
54.760      57.948
Status                             -2.0819      0.312     -6.665      0.000
-2.695     -1.469
Adult Mortality                     -0.0200      0.001    -20.484      0.000
-0.022     -0.018
=====
```

infant deaths	0.0945	0.010	9.703	0.000
0.075	0.114			
Alcohol	0.0403	0.031	1.288	0.198
-0.021	0.102			
Hepatitis B	-0.0209	0.005	-4.304	0.000
-0.030	-0.011			
BMI	0.0487	0.006	8.014	0.000
0.037	0.061			
under-five deaths	-0.0715	0.007	-10.003	0.000
-0.085	-0.057			
Polio	0.0272	0.005	4.963	0.000
0.016	0.038			
Total expenditure	0.0767	0.042	1.835	0.067
-0.005	0.159			
Diphtheria	0.0456	0.006	7.811	0.000
0.034	0.057			
HIV/AIDS	-0.4969	0.024	-20.599	0.000
-0.544	-0.450			
thinness 1-19 years	-0.0710	0.029	-2.427	0.015
-0.128	-0.014			
Income composition of resources	6.4837	0.768	8.437	0.000
4.977	7.991			
Schooling	0.6909	0.051	13.548	0.000
0.591	0.791			

Omnibus:	110.689	Durbin-Watson:	1.979
Prob(Omnibus):	0.000	Jarque-Bera (JB):	320.883
Skew:	-0.238	Prob(JB):	2.09e-70
Kurtosis:	4.876	Cond. No.	2.45e+03

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 2.45e+03. This might indicate that there are strong multicollinearity or other numerical problems.

```
[43]: # Create a dataframe that will contain the names of all the feature variables
      →and their respective VIFs
vif = pd.DataFrame()
vif['Features'] = X_train_rfe1.columns
vif['VIF'] = [variance_inflation_factor(X_train_rfe1.values, i) for i in
      →range(X_train_rfe1.shape[1])]
vif['VIF'] = round(vif['VIF'], 2)
vif = vif.sort_values(by = "VIF", ascending = False)
vif
```

```
[43]:
```

	Features	VIF
6	under-five deaths	177.82
2	infant deaths	177.16
13	Schooling	44.55
12	Income composition of resources	30.42
9	Diphtheria	30.30
7	Polio	26.28
4	Hepatitis B	18.99
5	BMI	8.19
8	Total expenditure	7.74
0	Status	7.10
1	Adult Mortality	4.41
3	Alcohol	4.35
11	thinness 1-19 years	4.07
10	HIV/AIDS	1.70

Since the variable under-five deaths is having a very high VIF score, we would remove the feature from training dataset

```
[44]: # Dropping insignificant variables

X_train_rfe2 = X_train_rfe1.drop('under-five deaths', 1,)

# Adding a constant variable and Build a second fitted model

X_train_rfe2c = sm.add_constant(X_train_rfe2)
lm_rfe2 = sm.OLS(y_train, X_train_rfe2c).fit()

#Summary of linear model
print(lm_rfe2.summary())
```

```

                                OLS Regression Results
=====
Dep. Variable:                Life expectancy    R-squared:                0.811
Model:                        OLS              Adj. R-squared:           0.810
Method:                      Least Squares     F-statistic:             675.4
Date:                        Sun, 20 Dec 2020   Prob (F-statistic):      0.00
Time:                        16:09:51          Log-Likelihood:          -5872.2
No. Observations:            2056             AIC:                    1.177e+04
Df Residuals:                2042             BIC:                    1.185e+04
Df Model:                    13
Covariance Type:             nonrobust
=====
=====
                                coef    std err          t      P>|t|
-----
[0.025    0.975]
-----
const                55.1520    0.823    67.005    0.000

```



53.538	56.766				
Status		-2.0413	0.320	-6.382	0.000
-2.668	-1.414				
Adult Mortality		-0.0204	0.001	-20.368	0.000
-0.022	-0.018				
infant deaths		-0.0025	0.001	-2.797	0.005
-0.004	-0.001				
Alcohol		-0.0038	0.032	-0.121	0.904
-0.066	0.058				
Hepatitis B		-0.0243	0.005	-4.887	0.000
-0.034	-0.015				
BMI		0.0502	0.006	8.071	0.000
0.038	0.062				
Polio		0.0307	0.006	5.484	0.000
0.020	0.042				
Total expenditure		0.0825	0.043	1.928	0.054
-0.001	0.166				
Diphtheria		0.0536	0.006	9.039	0.000
0.042	0.065				
HIV/AIDS		-0.5147	0.025	-20.893	0.000
-0.563	-0.466				
thinness 1-19 years		-0.0578	0.030	-1.932	0.053
-0.116	0.001				
Income composition of resources		7.1638	0.784	9.140	0.000
5.627	8.701				
Schooling		0.7020	0.052	13.448	0.000
0.600	0.804				
=====					
Omnibus:	110.170	Durbin-Watson:		1.988	
Prob(Omnibus):	0.000	Jarque-Bera (JB):		309.575	
Skew:	-0.250	Prob(JB):		5.98e-68	
Kurtosis:	4.834	Cond. No.		2.30e+03	
=====					

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 2.3e+03. This might indicate that there are strong multicollinearity or other numerical problems.

```
[45]: # Create a dataframe that will contain the names of all the feature variables
      ↪and their respective VIFs
vif = pd.DataFrame()
vif['Features'] = X_train_rfe2.columns
vif['VIF'] = [variance_inflation_factor(X_train_rfe2.values, i) for i in
      ↪range(X_train_rfe2.shape[1])]
vif['VIF'] = round(vif['VIF'], 2)
```

```
vif = vif.sort_values(by = "VIF", ascending = False)
vif
```

```
[45]:
```

	Features	VIF
12	Schooling	44.53
11	Income composition of resources	30.29
8	Diphtheria	29.80
6	Polio	26.22
4	Hepatitis B	18.76
5	BMI	8.19
7	Total expenditure	7.73
0	Status	7.06
1	Adult Mortality	4.38
3	Alcohol	4.23
10	thinness 1-19 years	4.07
9	HIV/AIDS	1.69
2	infant deaths	1.47

Since the variable Alcohol is having a very high p value, we would remove the feature from training dataset

```
[46]: # Dropping insignificant variables

X_train_rfe3 = X_train_rfe2.drop('Alcohol', 1,)

# Adding a constant variable and Build a second fitted model

X_train_rfe3c = sm.add_constant(X_train_rfe3)
lm_rfe3 = sm.OLS(y_train, X_train_rfe3c).fit()

#Summary of linear model
print(lm_rfe3.summary())
```

```

                                OLS Regression Results
=====
Dep. Variable:                Life expectancy    R-squared:                0.811
Model:                        OLS               Adj. R-squared:           0.810
Method:                       Least Squares      F-statistic:              732.0
Date:                         Sun, 20 Dec 2020    Prob (F-statistic):       0.00
Time:                         16:09:51          Log-Likelihood:           -5872.2
No. Observations:             2056              AIC:                     1.177e+04
Df Residuals:                 2043              BIC:                     1.184e+04
Df Model:                      12
Covariance Type:              nonrobust
=====
=====
                                coef    std err          t      P>|t|
-----
[0.025    0.975]
-----
```

```

-----
const                55.1373    0.814    67.751    0.000
53.541    56.733
Status              -2.0259    0.293    -6.904    0.000
-2.601    -1.450
Adult Mortality     -0.0204    0.001   -20.439    0.000
-0.022    -0.018
infant deaths       -0.0025    0.001    -2.826    0.005
-0.004    -0.001
Hepatitis B         -0.0243    0.005    -4.887    0.000
-0.034    -0.015
BMI                 0.0502    0.006     8.073    0.000
0.038     0.062
Polio               0.0307    0.006     5.484    0.000
0.020     0.042
Total expenditure   0.0819    0.042     1.928    0.054
-0.001     0.165
Diphtheria          0.0536    0.006     9.043    0.000
0.042     0.065
HIV/AIDS           -0.5149    0.025   -20.942    0.000
-0.563    -0.467
thinness 1-19 years -0.0571    0.029    -1.949    0.051
-0.114     0.000
Income composition of resources  7.1646    0.784     9.144    0.000
5.628     8.701
Schooling           0.7009    0.051    13.647    0.000
0.600     0.802
=====
Omnibus:              110.297    Durbin-Watson:              1.988
Prob(Omnibus):         0.000    Jarque-Bera (JB):          309.319
Skew:                 -0.251    Prob(JB):                  6.80e-68
Kurtosis:              4.833    Cond. No.                  2.28e+03
=====

```

Warnings:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 2.28e+03. This might indicate that there are strong multicollinearity or other numerical problems.

```

[47]: # Create a dataframe that will contain the names of all the feature variables
      ↪and their respective VIFs
vif = pd.DataFrame()
vif['Features'] = X_train_rfe3.columns
vif['VIF'] = [variance_inflation_factor(X_train_rfe3.values, i) for i in
      ↪range(X_train_rfe3.shape[1])]
vif['VIF'] = round(vif['VIF'], 2)

```

```
vif = vif.sort_values(by = "VIF", ascending = False)
vif
```

```
[47]:
```

	Features	VIF
11	Schooling	42.11
10	Income composition of resources	30.28
7	Diphtheria	29.80
5	Polio	26.16
3	Hepatitis B	18.73
4	BMI	8.18
6	Total expenditure	7.49
0	Status	6.05
1	Adult Mortality	4.30
9	thinness 1-19 years	3.96
8	HIV/AIDS	1.69
2	infant deaths	1.45

Since the variable **Schooling** is having a very high VIF score, we would remove the feature from training dataset

```
[48]: # Dropping insignificant variables

X_train_rfe4 = X_train_rfe3.drop('Schooling', 1,)

# Adding a constant variable and Build a second fitted model

X_train_rfe4c = sm.add_constant(X_train_rfe4)
lm_rfe4 = sm.OLS(y_train, X_train_rfe4c).fit()

#Summary of linear model
print(lm_rfe4.summary())
```

```

                                OLS Regression Results
=====
Dep. Variable:          Life expectancy      R-squared:                0.794
Model:                  OLS                 Adj. R-squared:           0.793
Method:                 Least Squares        F-statistic:              716.7
Date:                  Sun, 20 Dec 2020      Prob (F-statistic):       0.00
Time:                  16:09:52              Log-Likelihood:          -5961.9
No. Observations:      2056                 AIC:                     1.195e+04
Df Residuals:          2044                 BIC:                     1.202e+04
Df Model:              11
Covariance Type:       nonrobust
=====
=====
                                coef      std err          t      P>|t|
-----
[0.025      0.975]
-----
-----
```

const		58.7499	0.804	73.101	0.000
57.174	60.326				
Status		-2.6868	0.302	-8.889	0.000
-3.280	-2.094				
Adult Mortality		-0.0217	0.001	-20.948	0.000
-0.024	-0.020				
infant deaths		-0.0029	0.001	-3.092	0.002
-0.005	-0.001				
Hepatitis B		-0.0249	0.005	-4.801	0.000
-0.035	-0.015				
BMI		0.0626	0.006	9.740	0.000
0.050	0.075				
Polio		0.0360	0.006	6.188	0.000
0.025	0.047				
Total expenditure		0.1155	0.044	2.608	0.009
0.029	0.202				
Diphtheria		0.0572	0.006	9.249	0.000
0.045	0.069				
HIV/AIDS		-0.4918	0.026	-19.200	0.000
-0.542	-0.442				
thinness 1-19 years		-0.0830	0.031	-2.719	0.007
-0.143	-0.023				
Income composition of resources		13.9929	0.630	22.220	0.000
12.758	15.228				

Omnibus:	88.663	Durbin-Watson:	2.009
Prob(Omnibus):	0.000	Jarque-Bera (JB):	262.468
Skew:	-0.121	Prob(JB):	1.01e-57
Kurtosis:	4.734	Cond. No.	2.27e+03

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 2.27e+03. This might indicate that there are strong multicollinearity or other numerical problems.

```
[49]: # Create a dataframe that will contain the names of all the feature variables
      →and their respective VIFs
vif = pd.DataFrame()
vif['Features'] = X_train_rfe4.columns
vif['VIF'] = [variance_inflation_factor(X_train_rfe4.values, i) for i in
      →range(X_train_rfe4.shape[1])]
vif['VIF'] = round(vif['VIF'], 2)
vif = vif.sort_values(by = "VIF", ascending = False)
vif
```

[49]:

	Features	VIF
7	Diphtheria	29.69
5	Polio	25.80
3	Hepatitis B	18.59
10	Income composition of resources	13.90
4	BMI	7.83
6	Total expenditure	7.22
0	Status	6.05
1	Adult Mortality	4.30
9	thinness 1-19 years	3.95
8	HIV/AIDS	1.69
2	infant deaths	1.45

Since the variable Diphtheria is having a very high VIF score, we would remove the feature from training dataset

[50]: *# Dropping insignificant variables*

```

X_train_rfe5 = X_train_rfe4.drop('Diphtheria', 1,)

# Adding a constant variable and Build a second fitted model

X_train_rfe5c = sm.add_constant(X_train_rfe5)
lm_rfe5 = sm.OLS(y_train, X_train_rfe5c).fit()

#Summary of linear model
print(lm_rfe5.summary())

# Create a dataframe that will contain the names of all the feature variables,
→and their respective VIFs
vif = pd.DataFrame()
vif['Features'] = X_train_rfe5.columns
vif['VIF'] = [variance_inflation_factor(X_train_rfe5.values, i) for i in
→range(X_train_rfe5.shape[1])]
vif['VIF'] = round(vif['VIF'], 2)
vif = vif.sort_values(by = "VIF", ascending = False)
vif

```

#### OLS Regression Results

```

=====
Dep. Variable:      Life expectancy      R-squared:      0.785
Model:              OLS                  Adj. R-squared:  0.784
Method:             Least Squares        F-statistic:    748.8
Date:               Sun, 20 Dec 2020      Prob (F-statistic): 0.00
Time:               16:09:52              Log-Likelihood: -6004.1
No. Observations:   2056                  AIC:           1.203e+04
Df Residuals:       2045                  BIC:           1.209e+04
Df Model:           10
Covariance Type:    nonrobust

```

```

=====
=====
                                coef      std err          t      P>|t|
[0.025      0.975]
-----
const                        59.1919      0.819      72.302      0.000
57.586      60.797
Status                       -2.6778      0.308     -8.682      0.000
-3.283     -2.073
Adult Mortality             -0.0221      0.001    -20.928      0.000
-0.024     -0.020
infant deaths               -0.0031      0.001     -3.256      0.001
-0.005     -0.001
Hepatitis B                 -0.0102      0.005     -2.018      0.044
-0.020     -0.000
BMI                          0.0645      0.007      9.838      0.000
0.052      0.077
Polio                       0.0649      0.005     12.954      0.000
0.055      0.075
Total expenditure           0.1491      0.045      3.310      0.001
0.061      0.237
HIV/AIDS                   -0.4937      0.026    -18.889      0.000
-0.545     -0.442
thinness 1-19 years        -0.0823      0.031     -2.644      0.008
-0.143     -0.021
Income composition of resources 14.7697      0.637     23.190      0.000
13.521     16.019
=====
Omnibus:                    92.191  Durbin-Watson:                2.014
Prob(Omnibus):              0.000  Jarque-Bera (JB):            273.609
Skew:                      -0.141  Prob(JB):                    3.86e-60
Kurtosis:                   4.765  Cond. No.                     2.17e+03
=====

```

#### Warnings:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 2.17e+03. This might indicate that there are strong multicollinearity or other numerical problems.

```

[50]:
Features  VIF
5         Polio 17.90
3         Hepatitis B 16.46
9  Income composition of resources 13.42
4         BMI 7.82
6         Total expenditure 7.12

```

0	Status	6.04
1	Adult Mortality	4.29
8	thinness 1-19 years	3.95
7	HIV/AIDS	1.69
2	infant deaths	1.45

[51]: *# Dropping insignificant variables*

```
X_train_rfe6 = X_train_rfe5.drop('Polio', 1,)

# Adding a constant variable and Build a second fitted model

X_train_rfe6c = sm.add_constant(X_train_rfe6)
lm_rfe6 = sm.OLS(y_train, X_train_rfe6c).fit()

#Summary of linear model
print(lm_rfe6.summary())

# Create a dataframe that will contain the names of all the feature variables,
→and their respective VIFs
vif = pd.DataFrame()
vif['Features'] = X_train_rfe6.columns
vif['VIF'] = [variance_inflation_factor(X_train_rfe6.values, i) for i in
→range(X_train_rfe6.shape[1])]
vif['VIF'] = round(vif['VIF'], 2)
vif = vif.sort_values(by = "VIF", ascending = False)
vif
```

#### OLS Regression Results

```
=====
Dep. Variable:      Life expectancy      R-squared:      0.768
Model:              OLS      Adj. R-squared:      0.767
Method:             Least Squares      F-statistic:      752.1
Date:               Sun, 20 Dec 2020      Prob (F-statistic):      0.00
Time:               16:09:52      Log-Likelihood:      -6085.2
No. Observations:      2056      AIC:      1.219e+04
Df Residuals:         2046      BIC:      1.225e+04
Df Model:              9
Covariance Type:      nonrobust
=====
```

```
=====
                                coef      std err          t      P>|t|
-----
[0.025      0.975]
-----
const                                61.5258      0.831      74.081      0.000
59.897      63.155
Status                               -2.8026      0.321     -8.742      0.000
```



-3.431	-2.174				
Adult Mortality		-0.0232	0.001	-21.164	0.000
-0.025	-0.021				
infant deaths		-0.0038	0.001	-3.849	0.000
-0.006	-0.002				
Hepatitis B		0.0135	0.005	2.766	0.006
0.004	0.023				
BMI		0.0703	0.007	10.336	0.000
0.057	0.084				
Total expenditure		0.1839	0.047	3.932	0.000
0.092	0.276				
HIV/AIDS		-0.4923	0.027	-18.113	0.000
-0.546	-0.439				
thinness 1-19 years		-0.0744	0.032	-2.297	0.022
-0.138	-0.011				
Income composition of resources		16.2560	0.652	24.951	0.000
14.978	17.534				

Omnibus:	124.779	Durbin-Watson:	1.992
Prob(Omnibus):	0.000	Jarque-Bera (JB):	401.157
Skew:	-0.247	Prob(JB):	7.76e-88
Kurtosis:	5.107	Cond. No.	2.06e+03

#### Warnings:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 2.06e+03. This might indicate that there are strong multicollinearity or other numerical problems.

[51]:

	Features	VIF
3	Hepatitis B	12.86
8	Income composition of resources	11.81
4	BMI	7.70
5	Total expenditure	6.94
0	Status	6.00
1	Adult Mortality	4.29
7	thinness 1-19 years	3.92
6	HIV/AIDS	1.69
2	infant deaths	1.45

[52]: *# Dropping insignificant variables*

```
X_train_rfe7 = X_train_rfe6.drop('Hepatitis B', 1,)
```

*# Adding a constant variable and Build a second fitted model*

```
X_train_rfe7c = sm.add_constant(X_train_rfe7)
```

```

lm_rfe7 = sm.OLS(y_train, X_train_rfe7c).fit()

#Summary of linear model
print(lm_rfe7.summary())

# Create a dataframe that will contain the names of all the feature variables,
→and their respective VIFs
vif = pd.DataFrame()
vif['Features'] = X_train_rfe7.columns
vif['VIF'] = [variance_inflation_factor(X_train_rfe7.values, i) for i in
→range(X_train_rfe7.shape[1])]
vif['VIF'] = round(vif['VIF'], 2)
vif = vif.sort_values(by = "VIF", ascending = False)
vif

```

#### OLS Regression Results

```

=====
Dep. Variable:          Life expectancy    R-squared:                0.767
Model:                  OLS               Adj. R-squared:         0.766
Method:                 Least Squares     F-statistic:             842.4
Date:                  Sun, 20 Dec 2020   Prob (F-statistic):      0.00
Time:                  16:09:52          Log-Likelihood:         -6089.0
No. Observations:      2056             AIC:                   1.220e+04
Df Residuals:          2047             BIC:                   1.225e+04
Df Model:               8
Covariance Type:       nonrobust
=====

```

```

=====
                                coef    std err          t      P>|t|
-----
[0.025    0.975]
-----
const                                62.5312      0.748     83.601     0.000
61.064    63.998
Status                             -2.8044      0.321    -8.733     0.000
-3.434    -2.175
Adult Mortality                     -0.0233      0.001   -21.330     0.000
-0.025    -0.021
infant deaths                       -0.0043      0.001    -4.378     0.000
-0.006    -0.002
BMI                                0.0712      0.007    10.456     0.000
0.058     0.085
Total expenditure                   0.1859      0.047     3.971     0.000
0.094     0.278
HIV/AIDS                           -0.4957      0.027   -18.225     0.000
-0.549    -0.442
thinness 1-19 years                -0.0701      0.032    -2.164     0.031

```

```

-0.134      -0.007
Income composition of resources    16.3814      0.651      25.164      0.000
15.105      17.658
=====
Omnibus:                        123.339    Durbin-Watson:                1.995
Prob(Omnibus):                  0.000    Jarque-Bera (JB):                400.688
Skew:                          -0.237    Prob(JB):                        9.81e-88
Kurtosis:                      5.110    Cond. No.                        1.90e+03
=====

```

#### Warnings:

```

[1] Standard Errors assume that the covariance matrix of the errors is correctly
specified.
[2] The condition number is large, 1.9e+03. This might indicate that there are
strong multicollinearity or other numerical problems.

```

```

[52]:
Features    VIF
7  Income composition of resources  9.62
3                               BMI  7.49
4           Total expenditure  6.53
0                               Status  5.56
1           Adult Mortality  4.25
6           thinness 1-19 years  3.75
5                               HIV/AIDS  1.68
2           infant deaths  1.41

```

## 2.6 Approach 3 : Stepwise Regression

```
[53]: X_train.columns
```

```

[53]: Index(['Year', 'Status', 'Adult Mortality', 'infant deaths', 'Alcohol',
'percentage expenditure', 'Hepatitis B', 'Measles', 'BMI',
'under-five deaths', 'Polio', 'Total expenditure', 'Diphtheria',
'HIV/AIDS', 'GDP', 'Population', 'thinness 1-19 years',
'thinness 5-9 years', 'Income composition of resources', 'Schooling'],
dtype='object')

```

```
[54]: X_train.head()
```

```

[54]:
   Year  Status  Adult Mortality  infant deaths  Alcohol  \
270  2001      1           21.0             0      4.90
1687 2011      1          124.0            34      5.30
822  2011      1          197.0             2      2.37
2030 2008      1          217.0            60      4.21
363  2004      1           17.0            81      6.85

   percentage expenditure  Hepatitis B  Measles  BMI  under-five deaths  \
270          251.658693          96.0      0  41.4              0

```

1687	1117.196097	98.0	3	6.8	39
822	549.278308	89.0	0	53.4	2
2030	155.476762	88.0	341	21.6	78
363	186.609049	96.0	0	46.9	93

	Polio	Total expenditure	Diphtheria	HIV/AIDS	GDP \
270	96.0	4.50	96.0	0.4	3419.275719
1687	97.0	6.40	97.0	0.1	9834.472689
822	9.0	6.81	89.0	0.3	3736.587130
2030	91.0	4.50	91.0	0.1	1919.466195
363	99.0	7.70	99.0	0.1	3623.476670

	Population	thinness 1-19 years	thinness 5-9 years \
270	254984.0	3.7	3.7
1687	119917.0	1.6	1.6
822	619256.0	1.7	1.6
2030	9751864.0	1.0	9.7
363	184738458.0	3.2	3.2

	Income composition of resources	Schooling
270	0.677	11.8
1687	0.745	12.6
822	0.666	13.0
2030	0.655	11.5
363	0.695	14.0

```
[55]: X_train.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 2056 entries, 270 to 2863
Data columns (total 20 columns):
```

#	Column	Non-Null Count	Dtype
---	-----	-----	-----
0	Year	2056 non-null	int64
1	Status	2056 non-null	int64
2	Adult Mortality	2056 non-null	float64
3	infant deaths	2056 non-null	int64
4	Alcohol	2056 non-null	float64
5	percentage expenditure	2056 non-null	float64
6	Hepatitis B	2056 non-null	float64
7	Measles	2056 non-null	int64
8	BMI	2056 non-null	float64
9	under-five deaths	2056 non-null	int64
10	Polio	2056 non-null	float64
11	Total expenditure	2056 non-null	float64
12	Diphtheria	2056 non-null	float64
13	HIV/AIDS	2056 non-null	float64
14	GDP	2056 non-null	float64

```

15 Population                2056 non-null    float64
16 thinness 1-19 years       2056 non-null    float64
17 thinness 5-9 years        2056 non-null    float64
18 Income composition of resources 2056 non-null    float64
19 Schooling                 2056 non-null    float64
dtypes: float64(15), int64(5)
memory usage: 337.3 KB

```

```
[56]: y_train.head()
```

```

[56]:      Life expectancy
270          68.2
1687         76.1
822          72.0
2030         67.5
363          72.0

```

```
[57]: y_train.info()
```

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 2056 entries, 270 to 2863
Data columns (total 1 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Life expectancy  2056 non-null   float64
dtypes: float64(1)
memory usage: 32.1 KB

```

```
[ ]: ## By David Dale https://datascience.stackexchange.com/users/24162/david-dale
```

```

def stepwise_selection(X, y,
                      initial_list=[],
                      threshold_in=0.01,
                      threshold_out = 0.05,
                      verbose=True):
    """ Perform a forward-backward feature selection
    based on p-value from statsmodels.api.OLS
    Arguments:
        X - pandas.DataFrame with candidate features
        y - list-like with the target
        initial_list - list of features to start with (column names of X)
        threshold_in - include a feature if its p-value < threshold_in
        threshold_out - exclude a feature if its p-value > threshold_out
        verbose - whether to print the sequence of inclusions and exclusions
    Returns: list of selected features
    Always set threshold_in < threshold_out to avoid infinite looping.
    See https://en.wikipedia.org/wiki/Stepwise\_regression for the details
    """

```

```

included = list(initial_list)
while True:
    changed=False
    # forward step
    excluded = list(set(X.columns)-set(included))
    new_pval = pd.Series(index=excluded)
    for new_column in excluded:
        model = sm.OLS(y, sm.add_constant(pd.
→DataFrame(X[included+[new_column]]))).fit()
        new_pval[new_column] = model.pvalues[new_column]
    best_pval = new_pval.min()
    if best_pval < threshold_in:
        best_feature = new_pval.argmin()
        included.append(best_feature)
        changed=True
        if verbose:
            print('Add {:30} with p-value {:.6}'.format(best_feature,
→best_pval))

    # backward step
    # print(included)
    model = sm.OLS(y, sm.add_constant(pd.DataFrame(X[included]))).fit()
    # use all coefs except intercept
    pvalues = model.pvalues.iloc[1:]
    worst_pval = pvalues.max() # null if pvalues is empty
    if worst_pval > threshold_out:
        changed=True
        worst_feature = pvalues.argmax()
        included.remove(worst_feature)
        if verbose:
            print('Drop {:30} with p-value {:.6}'.format(worst_feature,
→worst_pval))
    if not changed:
        break
    return included

result = stepwise_selection(X_train, y_train)

print('resulting features:')
print(result)

```

[59]: X\_train\_stepwise = X\_train[['Schooling', 'Adult Mortality', 'HIV/AIDS',  
→'Diphtheria', 'BMI', 'Income composition of resources', 'Status',  
→'percentage expenditure', 'Polio', 'Measles', 'Hepatitis B', 'under-five'  
→deaths', 'infant deaths', 'thinness 1-19 years']]

*# Adding a constant variable and Build a second fitted model*

```
X_train_stepwise = sm.add_constant(X_train_stepwise)
lm_stepwise = sm.OLS(y_train, X_train_stepwise).fit()

#Summary of linear model
print(lm_stepwise.summary())
```

#### OLS Regression Results

```
=====
Dep. Variable:          Life expectancy    R-squared:                0.823
Model:                  OLS               Adj. R-squared:         0.822
Method:                 Least Squares     F-statistic:            677.3
Date:                   Sun, 20 Dec 2020   Prob (F-statistic):      0.00
Time:                   16:11:00          Log-Likelihood:         -5807.2
No. Observations:       2056             AIC:                   1.164e+04
Df Residuals:           2041             BIC:                   1.173e+04
Df Model:               14
Covariance Type:        nonrobust
=====
```

```
=====
                                coef    std err          t      P>|t|
-----
[0.025    0.975]
-----
const                        56.8039      0.756     75.122     0.000
55.321    58.287
Schooling                    0.6985      0.050     14.021     0.000
0.601     0.796
Adult Mortality              -0.0197      0.001    -20.365     0.000
-0.022    -0.018
HIV/AIDS                    -0.4943      0.024    -20.731     0.000
-0.541    -0.448
Diphtheria                   0.0457      0.006      7.901     0.000
0.034     0.057
BMI                          0.0491      0.006      8.134     0.000
0.037     0.061
Income composition of resources  5.8765      0.763      7.704     0.000
4.380     7.372
Status                       -1.8727      0.291     -6.433     0.000
-2.444    -1.302
percentage expenditure        0.0003    5.07e-05     5.555     0.000
0.000     0.000
Polio                        0.0272      0.005      5.004     0.000
0.017     0.038
Measles                     -2.367e-05  9.35e-06     -2.530     0.011
-4.2e-05  -5.32e-06
Hepatitis B                  -0.0194      0.005     -4.010     0.000
```

-0.029	-0.010				
under-five deaths		-0.0701	0.007	-9.950	0.000
-0.084	-0.056				
infant deaths		0.0940	0.010	9.797	0.000
0.075	0.113				
thinness 1-19 years		-0.0868	0.028	-3.056	0.002
-0.143	-0.031				

```

=====
Omnibus:                106.745    Durbin-Watson:                1.971
Prob(Omnibus):           0.000    Jarque-Bera (JB):            308.003
Skew:                    -0.225    Prob(JB):                    1.31e-67
Kurtosis:                4.842    Cond. No.                    1.05e+05
=====

```

Warnings:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 1.05e+05. This might indicate that there are strong multicollinearity or other numerical problems.

## 2.7 Model Prediction and Evaluation

```
[60]: # Predicting the price of training set.
X_test_stepwise = X_test[['Schooling', 'Adult Mortality', 'HIV/AIDS',
    → 'Diphtheria', 'BMI', 'Income composition of resources', 'Status',
    → 'percentage expenditure', 'Polio', 'Measles', 'Hepatitis B', 'under-five
    → deaths', 'infant deaths', 'thinness 1-19 years']]
X_test_stepwise = sm.add_constant(X_test_stepwise)
actual          = y_test["Life expectancy"]
prediction = lm_stepwise.predict(X_test_stepwise)
```

```
[61]: #Evaluation: MSE
model_mse = mean_squared_error(prediction, actual)
print(model_mse)
```

15.972714682412724

```
[62]: def mean_absolute_percentage_error(y_true, y_pred):
    y_true, y_pred = np.array(y_true), np.array(y_pred)
    return np.mean(np.abs((y_true - y_pred) / y_true)) * 100
```

```
[63]: mean_absolute_percentage_error(actual, prediction)
```

[63]: 4.558248666207725

```
[64]: # Check for Linearity
sns.scatterplot(y_test['Life expectancy'], prediction)
plt.title('Check for Linearity')
plt.xlabel('Actual value')
```

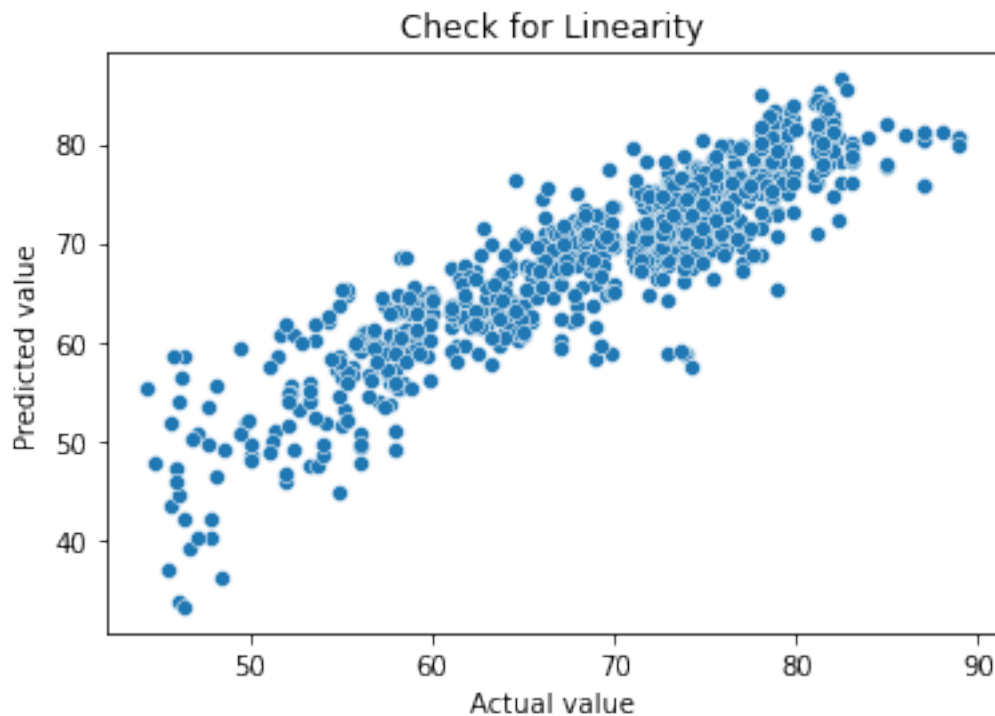


```
plt.ylabel('Predicted value')
```

```
/usr/local/lib/python3.6/dist-packages/seaborn/_decorators.py:43: FutureWarning:  
Pass the following variables as keyword args: x, y. From version 0.12, the only  
valid positional argument will be `data`, and passing other arguments without an  
explicit keyword will result in an error or misinterpretation.
```

```
FutureWarning
```

```
[64]: Text(0, 0.5, 'Predicted value')
```




```
[65]: # Plot the histogram of the error terms  
fig = plt.figure()  
sns.distplot((y_test['Life expectancy'] - prediction), bins = 20)  
fig.suptitle('Error Terms Analysis', fontsize = 20)  
plt.xlabel('Errors', fontsize = 18)
```

```
/usr/local/lib/python3.6/dist-packages/seaborn/distributions.py:2551:  
FutureWarning: `distplot` is a deprecated function and will be removed in a  
future version. Please adapt your code to use either `displot` (a figure-level  
function with similar flexibility) or `histplot` (an axes-level function for  
histograms).
```

```
warnings.warn(msg, FutureWarning)
```

```
[65]: Text(0.5, 0, 'Errors')
```



```
[ ]:
```